

## Text Classification 4395.001 Pranay Mantramurti Saige Wright

```
1 from google.colab import files
2
3 uploaded = files.upload()
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving PostPreQuestionsSmaller (1).csv to PostPreQuestionsSmaller (1).csv

```
1 import sklearn
2 import numpy as np
3 import pandas as pd
4 from sklearn.model_selection import train_test_split
5 import matplotlib.pyplot as plt
6
7 df = pd.read_csv('PostPreQuestionsSmaller (1).csv')
8
9 df.dropna(inplace = True)
10 df = df.astype({'label':'category'})
11 print(df[:5])
12 df['label'].value_counts()
13
14 #2 split train and test target: label,
15 Y_train, Y_test, X_train, X_test = train_test_split(df['label'],
16                                                    df['question'],
17                                                    test_size=0.2,
18                                                    random_state=1234)
19
```


	asin	question \
0	B017PICGL0	Are boots cold reliable?
1	B01HI8YEZS	is the main body rubber or plastic?
2	B07GZ1LF4R	Do you know the thread pitch/count for the tang?
3	B07BMNPRFN	The item ordered came tarnished and my son nee... these and now they look weird wit...

Saved successfully!

	item_name	hours_diff	label
0	Reebok Work Men's Zigkick RB7005 Work Shoe, Br...	-1	Pre
1	3C-Aone Galaxy S5 Case,Mangix Built-in Glass L...	21	Pre
2	ColdLand  14.00" Hand Forged Damascus Steel Bl...	166	Post
3	BEICHUANG Beidou 7 Stars Big Dipper Star Penda...	116	Post
4	Alla Lighting H8 H11 LED Bulbs Xtreme Super Br...	1021	Post

```
1 #creating bar chart distribution
2 import matplotlib.pyplot as plt
3 labels = df['label'].unique()
4 counts = []
5 for label in labels:
6     counts.append(len(df[df['label'] == label]))
7 plt.bar(labels,counts)
```

<BarContainer object of 2 artists>



The data set PrePostQuestions contains 232,492 product questions from 2019-2020, one per line. With columns question, asin, item\_name, hours\_diff, and label. The model should be able to predict the label of the questions based off the text. Our set that we have running through the code is the first 5000 questions from the original data set put in a smaller cvs file to read from

```
1 y_train_model = Y_train.copy()
2 y_test_model = Y_test.copy()
```

```
3 x_train_model = X_train.copy()
4 x_test_model = X_test.copy()
```

## ▼ Sequential Model

```
1 import tensorflow as tf
2 from tensorflow.keras.preprocessing.text import Tokenizer
3 from tensorflow.keras import layers, models
4
5 from sklearn.preprocessing import LabelEncoder
6 import pickle
7 import numpy as np
8 import pandas as pd
9
10 # set seed for reproducibility
11 np.random.seed(1234)
12
13 # set up X and Y
14 num_labels = 2
15 vocab_size = 25000
16 batch_size = 100
17
18 # fit the tokenizer on the training data
19 tokenizer = Tokenizer(num_words=vocab_size)
20 tokenizer.fit_on_texts(X_train)
21
22 x_train = tokenizer.texts_to_matrix(X_train, mode='tfidf')
23 x_test = tokenizer.texts_to_matrix(X_test, mode='tfidf')
```

```
1 encoder = LabelEncoder()
2 encoder.fit(Y_train)
3
4 temp_train = np.array(Y_train)
5 temp_test = np.array(Y_test)
6 y_train = encoder.transform(temp_train)
7
8 y_test = encoder.transform(temp_test)
```

```
1 # check shape
2 print("train shapes:", x_train.shape, y_train.shape)
3 print("test shapes:", x_test.shape, y_test.shape)
4 print("test first five labels:", y_test[:5])
```

```
train shapes: (3997, 25000) (3997,)
test shapes: (1000, 25000) (1000,)
test first five labels: [0 1 1 1 0]
```

```
1 model = models.Sequential()
2 model.add(layers.Dense(32, input_dim=vocab_size, kernel_initializer='normal', activation='relu'))
3 model.add(layers.Dense(1, kernel_initializer='normal', activation='sigmoid'))
4
5 model.compile(loss='binary_crossentropy',
6               optimizer='adam',
7               metrics=['accuracy'])
8
9 history = model.fit(x_train, y_train,
10                    batch_size=batch_size,
11                    epochs=10,
12                    verbose=1,
13                    validation_split=0.1)
```

```
Epoch 1/10
36/36 [=====] - 3s 13ms/step - loss: 0.6752 - accuracy: 0.5977 - val_loss: 0.6503 - val_accuracy: 0.6650
Epoch 2/10
36/36 [=====] - 0s 8ms/step - loss: 0.5646 - accuracy: 0.7843 - val_loss: 0.5824 - val_accuracy: 0.7100
Epoch 3/10
36/36 [=====] - 0s 8ms/step - loss: 0.4189 - accuracy: 0.8504 - val_loss: 0.5505 - val_accuracy: 0.7275
Epoch 4/10
36/36 [=====] - 0s 8ms/step - loss: 0.3042 - accuracy: 0.8832 - val_loss: 0.5610 - val_accuracy: 0.7300
Epoch 5/10
36/36 [=====] - 0s 8ms/step - loss: 0.2273 - accuracy: 0.9224 - val_loss: 0.5886 - val_accuracy: 0.7250
Epoch 6/10
36/36 [=====] - 0s 9ms/step - loss: 0.1757 - accuracy: 0.9458 - val_loss: 0.6254 - val_accuracy: 0.7150
Epoch 7/10
```

```

36/36 [=====] - 0s 8ms/step - loss: 0.1398 - accuracy: 0.9614 - val_loss: 0.6644 - val_accuracy: 0.6975
Epoch 8/10
36/36 [=====] - 0s 8ms/step - loss: 0.1134 - accuracy: 0.9716 - val_loss: 0.7040 - val_accuracy: 0.6925
Epoch 9/10
36/36 [=====] - 0s 8ms/step - loss: 0.0943 - accuracy: 0.9778 - val_loss: 0.7495 - val_accuracy: 0.6925
Epoch 10/10
36/36 [=====] - 0s 8ms/step - loss: 0.0798 - accuracy: 0.9811 - val_loss: 0.7837 - val_accuracy: 0.6925

1 # evaluate
2 score = model.evaluate(x_test, y_test, batch_size=batch_size, verbose=1)
3 print('Accuracy: ', score[1])

10/10 [=====] - 0s 5ms/step - loss: 0.8199 - accuracy: 0.6860
Accuracy: 0.685999895095825

1 pred = model.predict(x_test)
2 pred_labels = [1 if p>0.5 else 0 for p in pred]
3 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
4 print('accuracy score: ', accuracy_score(y_test, pred_labels))
5 print('precision score: ', precision_score(y_test, pred_labels))
6 print('recall score: ', recall_score(y_test, pred_labels))
7 print('f1 score: ', f1_score(y_test, pred_labels))

32/32 [=====] - 0s 3ms/step
accuracy score: 0.686
precision score: 0.7287066246056783
recall score: 0.7649006622516556
f1 score: 0.7463651050080775

```

## ▼ CNN

```

1 import tensorflow as tf
2 from tensorflow.keras import datasets, layers, models, preprocessing
3
4
5
6
7 x_train = [x.rjust(maxlen) for x in X_train]
8 x_test = [x.rjust(maxlen) for x in X_test]
9
10 temp_train = []
11 temp_test = []
12
13 for val in x_train:
14     temp = []
15     for letter in val:
16         temp.append(ord(letter))
17     temp_train.append(temp)
18
19 for val in x_test:
20     temp = []
21     for letter in val:
22         temp.append(ord(letter))
23     temp_test.append(temp)
24
25 x_train = np.array(temp_train)
26 x_test = np.array(temp_test)
27
28 temp_train = []
29 temp_test = []
30
31 for y in Y_train:
32     if y == 'Pre':
33         temp_train.append(0)
34     else:
35         temp_train.append(1)
36
37 for y in Y_test:
38     if y == 'Pre':
39         temp_test.append(0)
40     else:
41         temp_test.append(1)
42
43 y_train = np.array(temp_train)

```

Saved successfully!



```

44 y_test = np.array(temp_test)
45
46 model = models.Sequential()
47 model.add(layers.Embedding(max_features, 128, input_length=maxlen))
48 model.add(layers.Conv1D(32, 7, activation='relu'))
49 model.add(layers.MaxPooling1D(5))
50 model.add(layers.Conv1D(32, 7, activation='relu'))
51 model.add(layers.GlobalMaxPooling1D())
52 model.add(layers.Dense(1))
53
54 model.compile(optimizer=tf.keras.optimizers.Adamax(learning_rate=1e-3), # set learning rate
55               loss='binary_crossentropy',
56               metrics=['accuracy'])
57
58 history = model.fit(x_train, y_train, epochs=20, batch_size=128, validation_split=0.2)
59
60 from sklearn.metrics import classification_report
61
62 pred = model.predict(x_test)
63 pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
64 predict = []
65
66 for x in pred:
67     if x == 0:
68         predict.append('Pre')
69     else:
70         predict.append('Post')
71
72 print(classification_report(Y_test, predict))

```

Epoch 1/20  
25/25 [=====] - 1s 15ms/step - loss: 0.7228 - accuracy: 0.5931 - val\_loss: 0.6573 - val\_accuracy: 0.6450  
Epoch 2/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6599 - accuracy: 0.5906 - val\_loss: 0.6377 - val\_accuracy: 0.6162  
Epoch 3/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6513 - accuracy: 0.6203 - val\_loss: 0.6326 - val\_accuracy: 0.6413  
Epoch 4/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6439 - accuracy: 0.6700 - val\_loss: 0.6277 - val\_accuracy: 0.6737  
Epoch 5/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6370 - accuracy: 0.6891 - val\_loss: 0.6228 - val\_accuracy: 0.6825  
Epoch 6/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6302 - accuracy: 0.7022 - val\_loss: 0.6168 - val\_accuracy: 0.6787  
Epoch 7/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6240 - accuracy: 0.7091 - val\_loss: 0.6128 - val\_accuracy: 0.6938  
Epoch 8/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6171 - accuracy: 0.7157 - val\_loss: 0.6095 - val\_accuracy: 0.6988  
Epoch 9/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6107 - accuracy: 0.7122 - val\_loss: 0.6029 - val\_accuracy: 0.6875  
Epoch 10/20  
25/25 [=====] - 0s 6ms/step - loss: 0.6033 - accuracy: 0.7247 - val\_loss: 0.6009 - val\_accuracy: 0.7000  
Epoch 11/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5957 - accuracy: 0.7235 - val\_loss: 0.5954 - val\_accuracy: 0.7000  
Epoch 12/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5881 - accuracy: 0.7322 - val\_loss: 0.5924 - val\_accuracy: 0.7025  
Epoch 13/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5789 - accuracy: 0.7329 - val\_loss: 0.5844 - val\_accuracy: 0.7000  
Epoch 14/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5707 - accuracy: 0.7407 - val\_loss: 0.5817 - val\_accuracy: 0.6963  
Epoch 15/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5614 - accuracy: 0.7419 - val\_loss: 0.5771 - val\_accuracy: 0.7025  
Epoch 16/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5520 - accuracy: 0.7444 - val\_loss: 0.5707 - val\_accuracy: 0.7150  
Epoch 17/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5439 - accuracy: 0.7491 - val\_loss: 0.5682 - val\_accuracy: 0.7113  
Epoch 18/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5342 - accuracy: 0.7532 - val\_loss: 0.5680 - val\_accuracy: 0.7050  
Epoch 19/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5274 - accuracy: 0.7538 - val\_loss: 0.5615 - val\_accuracy: 0.7163  
Epoch 20/20  
25/25 [=====] - 0s 6ms/step - loss: 0.5158 - accuracy: 0.7635 - val\_loss: 0.5599 - val\_accuracy: 0.7100  
32/32 [=====] - 0s 2ms/step

	precision	recall	f1-score	support
Post	0.66	0.53	0.59	396
Pre	0.73	0.82	0.77	604
accuracy			0.71	1000
macro avg	0.69	0.68	0.68	1000
weighted avg	0.70	0.71	0.70	1000

## ▼ Embeddings

```

1 import tensorflow as tf
2 from tensorflow.keras import datasets, layers, models, preprocessing
3
4 max_features = 10000
5 maxlen = 500
6
7 x_train = [x.rjust(maxlen) for x in x_train_model]
8 x_test = [x.rjust(maxlen) for x in x_test_model]
9
10 temp_train = []
11 temp_test = []
12
13 for val in x_train:
14     temp = []
15     for letter in val:
16         temp.append(ord(letter))
17     temp_train.append(temp)
18
19 for val in x_test:
20     temp = []
21     for letter in val:
22         temp.append(ord(letter))
23     temp_test.append(temp)
24
25 x_train = np.array(temp_train)
26 x_test = np.array(temp_test)
27
28 temp_train = []
29 temp_test = []
30
31 for y in y_train_model:
32     if y == 'Pre':
33         temp_train.append(1)
34
35 for y in Y_test:
36     if y == 'Pre':
37         temp_test.append(0)
38     else:
39         temp_test.append(1)
40
41 y_train = np.array(temp_train)
42 y_test = np.array(temp_test)

```

Saved successfully!



## ▼ Default Embedding

```

1 #different embedding approaches - Default Approach
2
3 model = models.Sequential()
4 model.add(layers.Embedding(max_features, 128, input_length=maxlen))
5 model.add(layers.Conv1D(32, 7, activation='relu'))
6 model.add(layers.MaxPooling1D(5))
7 model.add(layers.Conv1D(32, 7, activation='relu'))
8 model.add(layers.GlobalMaxPooling1D())
9 model.add(layers.Dense(1))
10
11 model.compile(optimizer=tf.keras.optimizers.Adamax(learning_rate=1e-3), # set learning rate
12               loss='binary_crossentropy',
13               metrics=['accuracy'])
14
15 history = model.fit(x_train, y_train, epochs=20, batch_size=128, validation_split=0.2)

```

Epoch 1/20  
 25/25 [=====] - 1s 13ms/step - loss: 0.8395 - accuracy: 0.5812 - val\_loss: 0.6636 - val\_accuracy: 0.6137  
 Epoch 2/20  
 25/25 [=====] - 0s 6ms/step - loss: 0.6797 - accuracy: 0.5812 - val\_loss: 0.6563 - val\_accuracy: 0.6137  
 Epoch 3/20  
 25/25 [=====] - 0s 6ms/step - loss: 0.6707 - accuracy: 0.5812 - val\_loss: 0.6548 - val\_accuracy: 0.6137  
 Epoch 4/20

```

25/25 [=====] - 0s 7ms/step - loss: 0.6673 - accuracy: 0.5809 - val_loss: 0.6513 - val_accuracy: 0.6137
Epoch 5/20
25/25 [=====] - 0s 6ms/step - loss: 0.6640 - accuracy: 0.5821 - val_loss: 0.6493 - val_accuracy: 0.6175
Epoch 6/20
25/25 [=====] - 0s 6ms/step - loss: 0.6610 - accuracy: 0.5830 - val_loss: 0.6461 - val_accuracy: 0.6162
Epoch 7/20
25/25 [=====] - 0s 6ms/step - loss: 0.6582 - accuracy: 0.5862 - val_loss: 0.6438 - val_accuracy: 0.6187
Epoch 8/20
25/25 [=====] - 0s 6ms/step - loss: 0.6554 - accuracy: 0.5902 - val_loss: 0.6417 - val_accuracy: 0.6237
Epoch 9/20
25/25 [=====] - 0s 6ms/step - loss: 0.6527 - accuracy: 0.6046 - val_loss: 0.6398 - val_accuracy: 0.6300
Epoch 10/20
25/25 [=====] - 0s 6ms/step - loss: 0.6502 - accuracy: 0.6178 - val_loss: 0.6374 - val_accuracy: 0.6400
Epoch 11/20
25/25 [=====] - 0s 6ms/step - loss: 0.6476 - accuracy: 0.6300 - val_loss: 0.6358 - val_accuracy: 0.6413
Epoch 12/20
25/25 [=====] - 0s 6ms/step - loss: 0.6452 - accuracy: 0.6509 - val_loss: 0.6339 - val_accuracy: 0.6587
Epoch 13/20
25/25 [=====] - 0s 6ms/step - loss: 0.6427 - accuracy: 0.6475 - val_loss: 0.6323 - val_accuracy: 0.6625
Epoch 14/20
25/25 [=====] - 0s 6ms/step - loss: 0.6402 - accuracy: 0.6734 - val_loss: 0.6310 - val_accuracy: 0.6662
Epoch 15/20
25/25 [=====] - 0s 6ms/step - loss: 0.6378 - accuracy: 0.6678 - val_loss: 0.6280 - val_accuracy: 0.6600
Epoch 16/20
25/25 [=====] - 0s 7ms/step - loss: 0.6354 - accuracy: 0.6794 - val_loss: 0.6282 - val_accuracy: 0.6737
Epoch 17/20
25/25 [=====] - 0s 6ms/step - loss: 0.6330 - accuracy: 0.6822 - val_loss: 0.6256 - val_accuracy: 0.6700
Epoch 18/20
25/25 [=====] - 0s 6ms/step - loss: 0.6300 - accuracy: 0.7003 - val_loss: 0.6250 - val_accuracy: 0.6687
Epoch 19/20
25/25 [=====] - 0s 6ms/step - loss: 0.6269 - accuracy: 0.7016 - val_loss: 0.6208 - val_accuracy: 0.6662
Epoch 20/20
25/25 [=====] - 0s 6ms/step - loss: 0.6236 - accuracy: 0.7019 - val_loss: 0.6199 - val_accuracy: 0.6625

```

## ▼ Using Custom Embedding Layer, similar to Github example

```
1 from tensorflow.keras import layers
```

Saved successfully!

```

4 EMBEDDING_DIM = 128
5 MAX_SEQUENCE_LENGTH = 500
6
7 embedding_layer = layers.Embedding(max_features + 1,
8                                   EMBEDDING_DIM,
9                                   input_length=MAX_SEQUENCE_LENGTH)
10 int_sequences_input = keras.Input(shape=(None,), dtype="int64")
11 embedded_sequences = embedding_layer(int_sequences_input)
12 x = layers.Conv1D(128, 5, activation="relu")(embedded_sequences)
13 x = layers.MaxPooling1D(5)(x)
14 x = layers.Conv1D(128, 5, activation="relu")(x)
15 x = layers.MaxPooling1D(5)(x)
16 x = layers.Conv1D(128, 5, activation="relu")(x)
17 x = layers.GlobalMaxPooling1D()(x)
18 x = layers.Dense(128, activation="relu")(x)
19 x = layers.Dropout(0.5)(x)
20 preds = layers.Dense(1, activation="softmax")(x)
21 model = keras.Model(int_sequences_input, preds)
22
23 model.compile(
24     loss="binary_crossentropy", optimizer="rmsprop", metrics=["acc"]
25 )
26 model.fit(x_train, y_train, batch_size=128, epochs=20, validation_split = 0.2)

Epoch 1/20
25/25 [=====] - 12s 25ms/step - loss: 0.6745 - acc: 0.4188 - val_loss: 0.6413 - val_acc: 0.3862
Epoch 2/20
25/25 [=====] - 0s 8ms/step - loss: 0.6518 - acc: 0.4188 - val_loss: 0.6226 - val_acc: 0.3862
Epoch 3/20
25/25 [=====] - 0s 8ms/step - loss: 0.6203 - acc: 0.4188 - val_loss: 0.5875 - val_acc: 0.3862
Epoch 4/20
25/25 [=====] - 0s 8ms/step - loss: 0.5973 - acc: 0.4188 - val_loss: 0.6674 - val_acc: 0.3862
Epoch 5/20
25/25 [=====] - 0s 8ms/step - loss: 0.5859 - acc: 0.4188 - val_loss: 0.5926 - val_acc: 0.3862
Epoch 6/20
25/25 [=====] - 0s 8ms/step - loss: 0.5632 - acc: 0.4188 - val_loss: 0.5688 - val_acc: 0.3862
Epoch 7/20
25/25 [=====] - 0s 8ms/step - loss: 0.5529 - acc: 0.4188 - val_loss: 0.5756 - val_acc: 0.3862
Epoch 8/20

```

```

25/25 [=====] - 0s 8ms/step - loss: 0.5282 - acc: 0.4188 - val_loss: 0.5673 - val_acc: 0.3862
Epoch 9/20
25/25 [=====] - 0s 8ms/step - loss: 0.5209 - acc: 0.4188 - val_loss: 0.5650 - val_acc: 0.3862
Epoch 10/20
25/25 [=====] - 0s 8ms/step - loss: 0.4940 - acc: 0.4188 - val_loss: 0.6071 - val_acc: 0.3862
Epoch 11/20
25/25 [=====] - 0s 8ms/step - loss: 0.4812 - acc: 0.4188 - val_loss: 0.7809 - val_acc: 0.3862
Epoch 12/20
25/25 [=====] - 0s 8ms/step - loss: 0.4680 - acc: 0.4188 - val_loss: 0.6040 - val_acc: 0.3862
Epoch 13/20
25/25 [=====] - 0s 8ms/step - loss: 0.4474 - acc: 0.4188 - val_loss: 0.5822 - val_acc: 0.3862
Epoch 14/20
25/25 [=====] - 0s 8ms/step - loss: 0.4150 - acc: 0.4188 - val_loss: 0.6710 - val_acc: 0.3862
Epoch 15/20
25/25 [=====] - 0s 8ms/step - loss: 0.3929 - acc: 0.4188 - val_loss: 0.6136 - val_acc: 0.3862
Epoch 16/20
25/25 [=====] - 0s 8ms/step - loss: 0.3482 - acc: 0.4188 - val_loss: 0.7086 - val_acc: 0.3862
Epoch 17/20
25/25 [=====] - 0s 8ms/step - loss: 0.3198 - acc: 0.4188 - val_loss: 0.7541 - val_acc: 0.3862
Epoch 18/20
25/25 [=====] - 0s 9ms/step - loss: 0.3152 - acc: 0.4188 - val_loss: 0.7135 - val_acc: 0.3862
Epoch 19/20
25/25 [=====] - 0s 8ms/step - loss: 0.2601 - acc: 0.4188 - val_loss: 0.7399 - val_acc: 0.3862
Epoch 20/20
25/25 [=====] - 0s 8ms/step - loss: 0.2476 - acc: 0.4188 - val_loss: 0.7168 - val_acc: 0.3862
<keras.callbacks.History at 0x7f8c0f3dd070>

```

```

1 from sklearn.metrics import classification_report
2
3 pred = model.predict(x_test)
4 pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
5 predict = []
6
7 for x in pred:
8     if x == 0:
9         predict.append('Pre')
10    else:
11        predict.append('Post')
--

```

Saved successfully!



est\_model, predict))

```

32/32 [=====] - 0s 3ms/step
          precision    recall  f1-score   support

     Post       0.40        1.00        0.57        396
     Pre        0.00        0.00        0.00        604

 accuracy          0.20
 macro avg          0.20
 weighted avg       0.16

```

```

/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are :
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are :
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are :
_warn_prf(average, modifier, msg_start, len(result))

```

Analysis: Initially our data set was over 230,000 product questions and our models took a long time to train, so we cut the dataset down to about 5,000 product questions. The first sequential model we had an accuracy score of 68% which was surprising but moreover our precision was higher at 78%. We took a couple different approaches trying different architecture like CNN, and preprocessing our data different; encoding or values, different activation functions, and embedding matrices. After attempting alternate embedding formats, it ended up being incompatible with our dataset, or didn't make much of a difference, overall the CNN was best.



Saved successfully! 